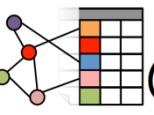


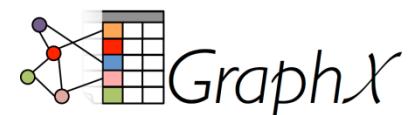
# Introduction to Spark

#5  Graph Analytics on  
Spark

Srini Penchikala  
[srinipenchikala@gmail.com](mailto:srinipenchikala@gmail.com)  
08/06/14

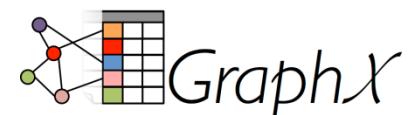
# About Me

- Solutions Architect
- Lead Editor for NoSQL at InfoQ.com
- Started as a Java Developer
- Current Focus: NoSQL Databases
  - Cassandra
  - MongoDB
  - Neo4j
  - Hazelcast
- Not a Data Scientist



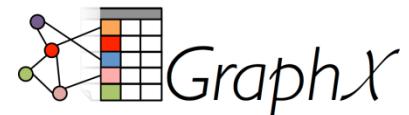
# Sources for this presentation

- Graph Analytics with GraphX\*\* (<https://databricks-training.s3.amazonaws.com/graph-analytics-with-graphx.html>)
- GraphX Programming Guide (<http://spark.apache.org/docs/latest/graphx-programming-guide.html>)
- Spark Summit 2014 Presentations
- AmpCamp 2014 Presentations



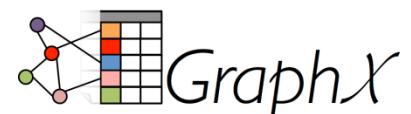
# Software Used for Today's Demos

- Spark 1.0.1
- JDK 1.6  
(<http://www.oracle.com/technetwork/java/javase/downloads/index-jsp-138363.html>)
- Scala IDE in Eclipse
- Windows

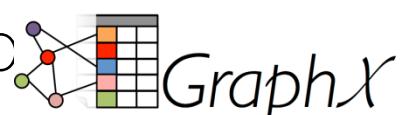


# How to Install Spark

- Spark 1.0.1  
(<https://spark.apache.org/downloads.html>)
- Cloudera VM  
([http://www.cloudera.com/content/support/en/downloads/quickstart\\_vms/cdh-5-1-x.html](http://www.cloudera.com/content/support/en/downloads/quickstart_vms/cdh-5-1-x.html))
- Vagrant/VirtualBox/Chef
- Others?

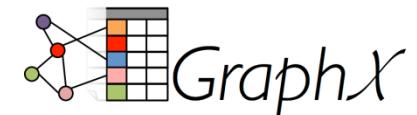


# How to run Spark

- Four ways to run Spark
  - Run Spark locally with one worker thread(no parallelism)
    - spark-shell.cmd or spark-shell.cmd local
  - Run locally with K worker threads (ideally set to # of cores)
    - spark-shell.cmd local[2]
  - Connect to a Spark standalone cluster
    - spark-shell.cmd spark://HOST:PORT
    - Example: MASTER=spark://localhost:7077
  - Connect to a Mesos cluster
    - spark-shell.cmd mesos://HOST:PC
- 
- The GraphX logo is located in the bottom right corner. It features a small network graph with four nodes connected by lines, each node containing a colored square (purple, orange, blue, green). To the right of the graph, the word "GraphX" is written in a stylized, lowercase font.

# Agenda

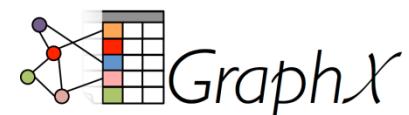
- Spark
- GraphX
- Use Cases
- Demos
- Q&A



# Spark

# Spark Background

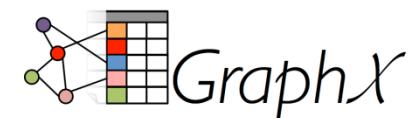
- Distributed Systems are the new standard (specifically, Data Parallel systems)
- Spark: An SDK for Big Data Applications
- Unified System With Libraries to Build a Complete Solution
- Full-featured Programming Environment
- Single, Consistent Interface for Developers to Write Against
- Runtimes available on several platforms



## **A Brief History:** Spark

Some key points about Spark:

- handles batch, interactive, and real-time within a single framework
- native integration with Java, Python, Scala
- programming at a higher level of abstraction
- more general: map/reduce is just one set of supported constructs



Info E Y M D R G S RW

https://lax6.readytalk.com/interface/flashView.jsp?url=services/lax6/core&uid36=ag3olj&simple=true&mwid=0&flashSk=f4bywt&flashDL=pond0-core&signed=null&jvm=Fla

Hide Panel Raise Hand Full Screen ReadyTalk

Audio Information  
Connected  
To listen via phone, dial:

Chairperson: Welcome everyone. We'll start momentarily. Please feel free to submit your questions for the Q&A via the chat.

Srini Penchikala: Please discuss how Spark is better than Hadoop for big data processing and analytics.

Srini Penchikala: Also, what are the limitations of Spark?

Chairperson: Reminder that you may enter questions for the Q&A in the Chat box at any time.

Chat with Presenter:  
<Type Message Here> Send

## What is Apache Spark?

- **Apache Spark** is a fast, general engine for large-scale data processing and analysis
  - Open source, developed at UC Berkeley
- **Written in Scala**
  - Functional programming language that runs in a JVM
- **Key Concepts**
  - Avoid the data bottleneck by distributing data when it is stored
  - Bring the processing to the data
  - Data stored in memory



cloudera Ask Bigger Questions



Info E Y M E R × 8 IW

https://lax6.readytalk.com/interface/flashView.jsp?url=services/lax6/core&uid36=ag3olj&simple=true&mwid=0&flashSk=f4bywt&flashDL=pond0-core&signed=null&jvm=Fla

Hide Panel Raise Hand Full Screen ReadyTalk

Audio Information  
Connected  
To listen via phone, dial:

**Chairperson:** Welcome everyone. We'll start momentarily. Please feel free to submit your questions for the Q&A via the chat.

**Srini Penchikala:** Please discuss how Spark is better than Hadoop for big data processing and analytics.

**Srini Penchikala:** Also, what are the limitations of Spark?

**Chairperson:** Reminder that you may enter questions for the Q&A in the Chat box at any time.

Chat with Presenter:  
<Type Message Here> Send

## Distributed Processing with the Spark Framework

```
graph TD; API[API  
Spark] <--> CC[Cluster Computing  
• Spark Standalone  
• YARN  
• Mesos]; API <--> Storage[Storage  
HDFS  
(Hadoop Distributed File System)]; CC <--> Storage;
```

The screenshot shows a web browser window with a presentation slide titled "RDD (Resilient Distributed Dataset)".

**Header:**

- Address bar: <https://lax6.readytalk.com/interface/flashView.jsp?uri=services/lax6/core&uid36=ag3olj&simple=true&mwid=0&flashSk=f4bywt&flashDL=pond0-core&signed=null&jvm=Fla>
- Toolbar: Standard browser icons for back, forward, search, etc.
- Buttons: Hide Panel, Raise Hand, Full Screen, ReadyTalk

**Left Panel (Audio Information):**

- Connected
- To listen via phone, dial:
- Chairperson: Welcome everyone. We'll start momentarily. Please feel free to submit your questions for the Q&A via the chat.
- Srini Penchikala: Please discuss how Spark is better than Hadoop for big data processing and analytics.
- Srini Penchikala: Also, what are the limitations of Spark?
- Chairperson: Reminder that you may enter questions for the Q&A in the Chat box at any time.

**Right Panel (Chat with Presenter):**

- Chat with Presenter:
- <Type Message Here>
- Send

**Content Area:**

## RDD (Resilient Distributed Dataset)

- RDD (Resilient Distributed Dataset)**
  - Resilient – if data in memory is lost, it can be recreated
  - Distributed – stored in memory across the cluster
  - Dataset – initial data can come from a file or created programmatically
- RDDs are the fundamental unit of data in Spark**
- Most of Spark programming is performing operations on RDDs**

**RDD Data Visualization:**

RDD
data
data
data
data...

**Footer:**

cloudera Ask Bigger Questions

# Spark Execution Model

1. Create DAG of RDDs to represent computation
2. Create logical execution plan for DAG
3. Schedule and execute individual tasks

## From Challenges to Solutions

Challenges	Solutions
Clusters hard to set up and manage	Hosted platform
Need to integrate a zoo of tools	Apache Spark
Tools are hard to use	Interactive Workspace

# GraphX

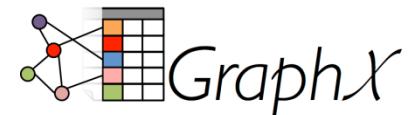
# What is GraphX

- New Spark API for:
  - Graphs (e.g., Web-Graphs and Social Networks)
  - Graph-parallel computation (e.g., PageRank and Collaborative Filtering)
- Extends Spark RDD abstraction by introducing Resilient Distributed Property Graph
- Property Graph: a directed multi-graph with properties attached to each vertex and edge
- Set of operators (e.g., subgraph, joinVertices, & mapReduceTriplets)
- Collection of graph algorithms
- Builders to simplify graph analytics tasks



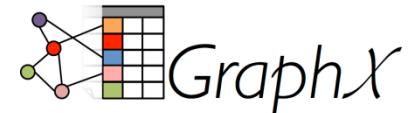
# What is GraphX - 2

- Currently only available in Scala
- Java and Python bindings (future)
- ToDo: Link to GraphX API



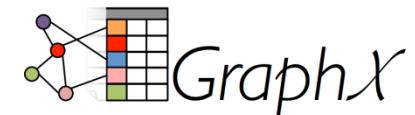
# Property Graphs

- Like RDDs, property graphs are immutable, distributed, and fault-tolerant.

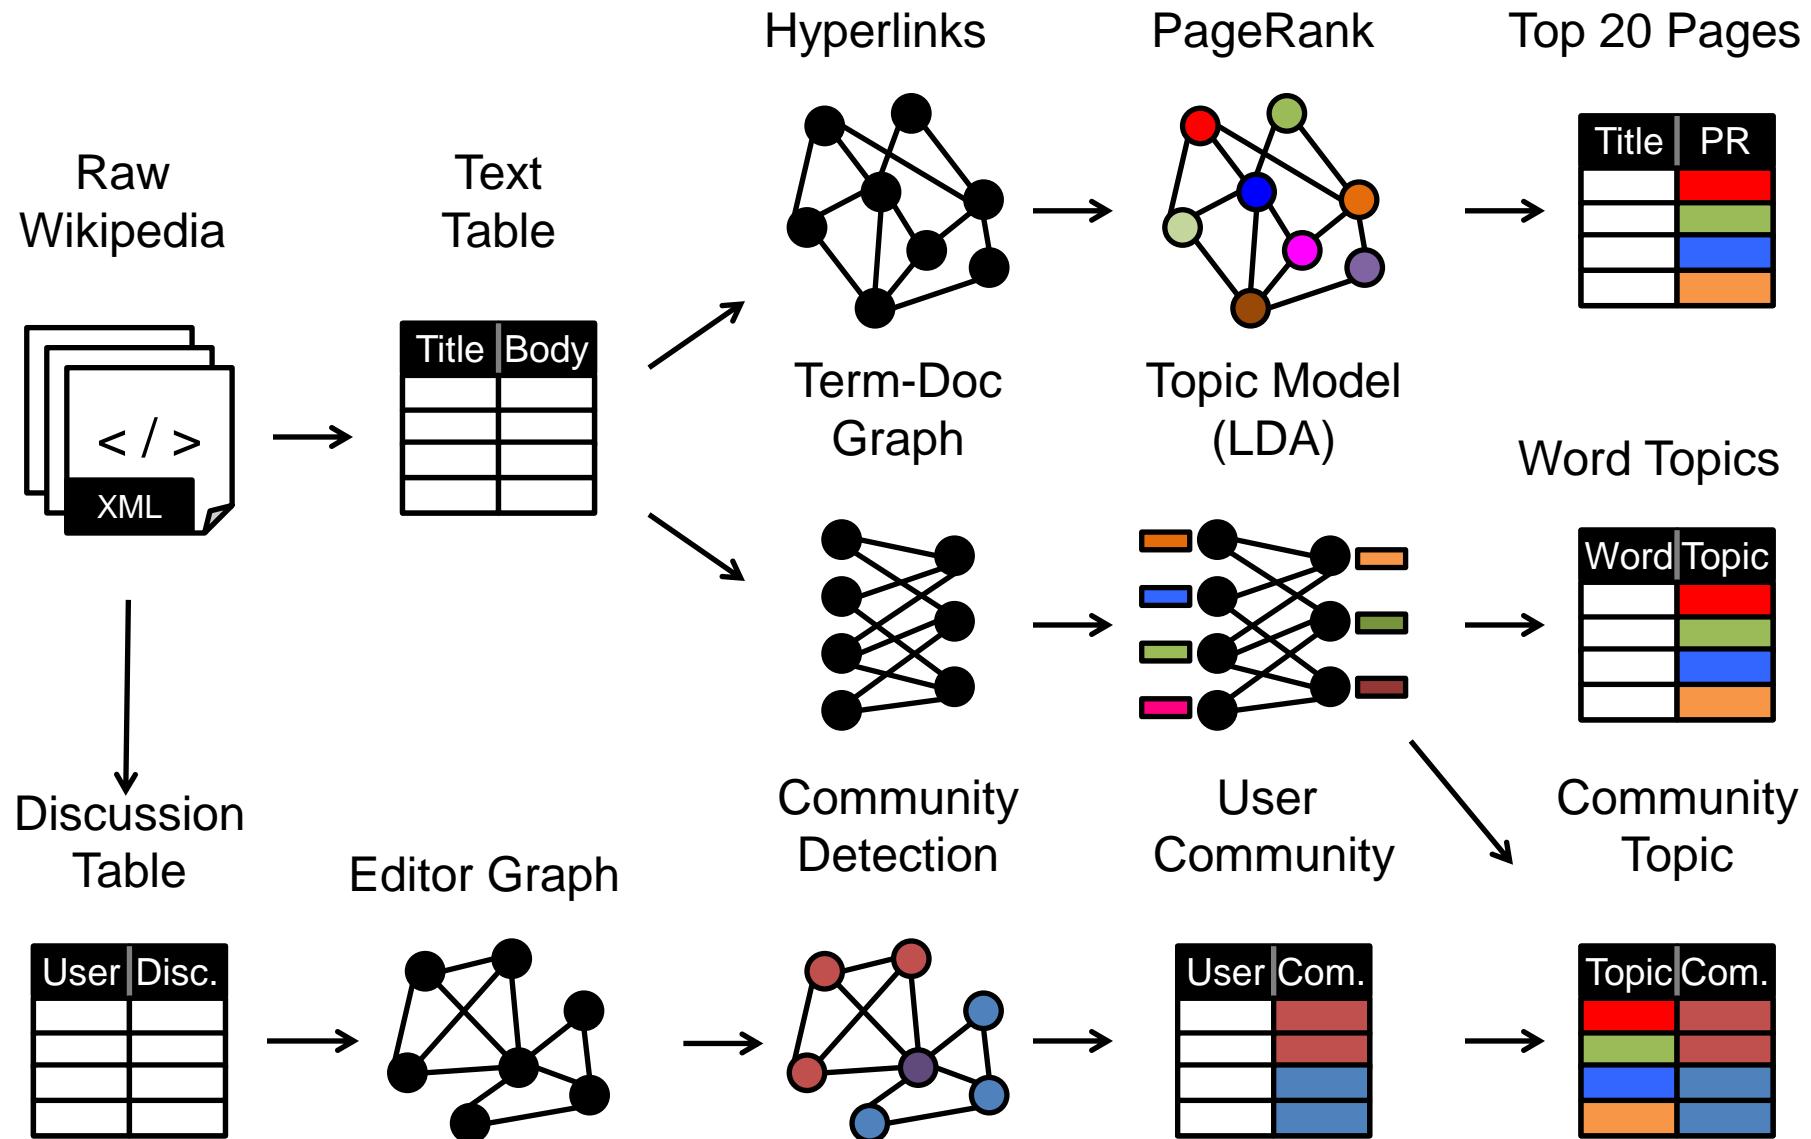


# Graph Algorithms

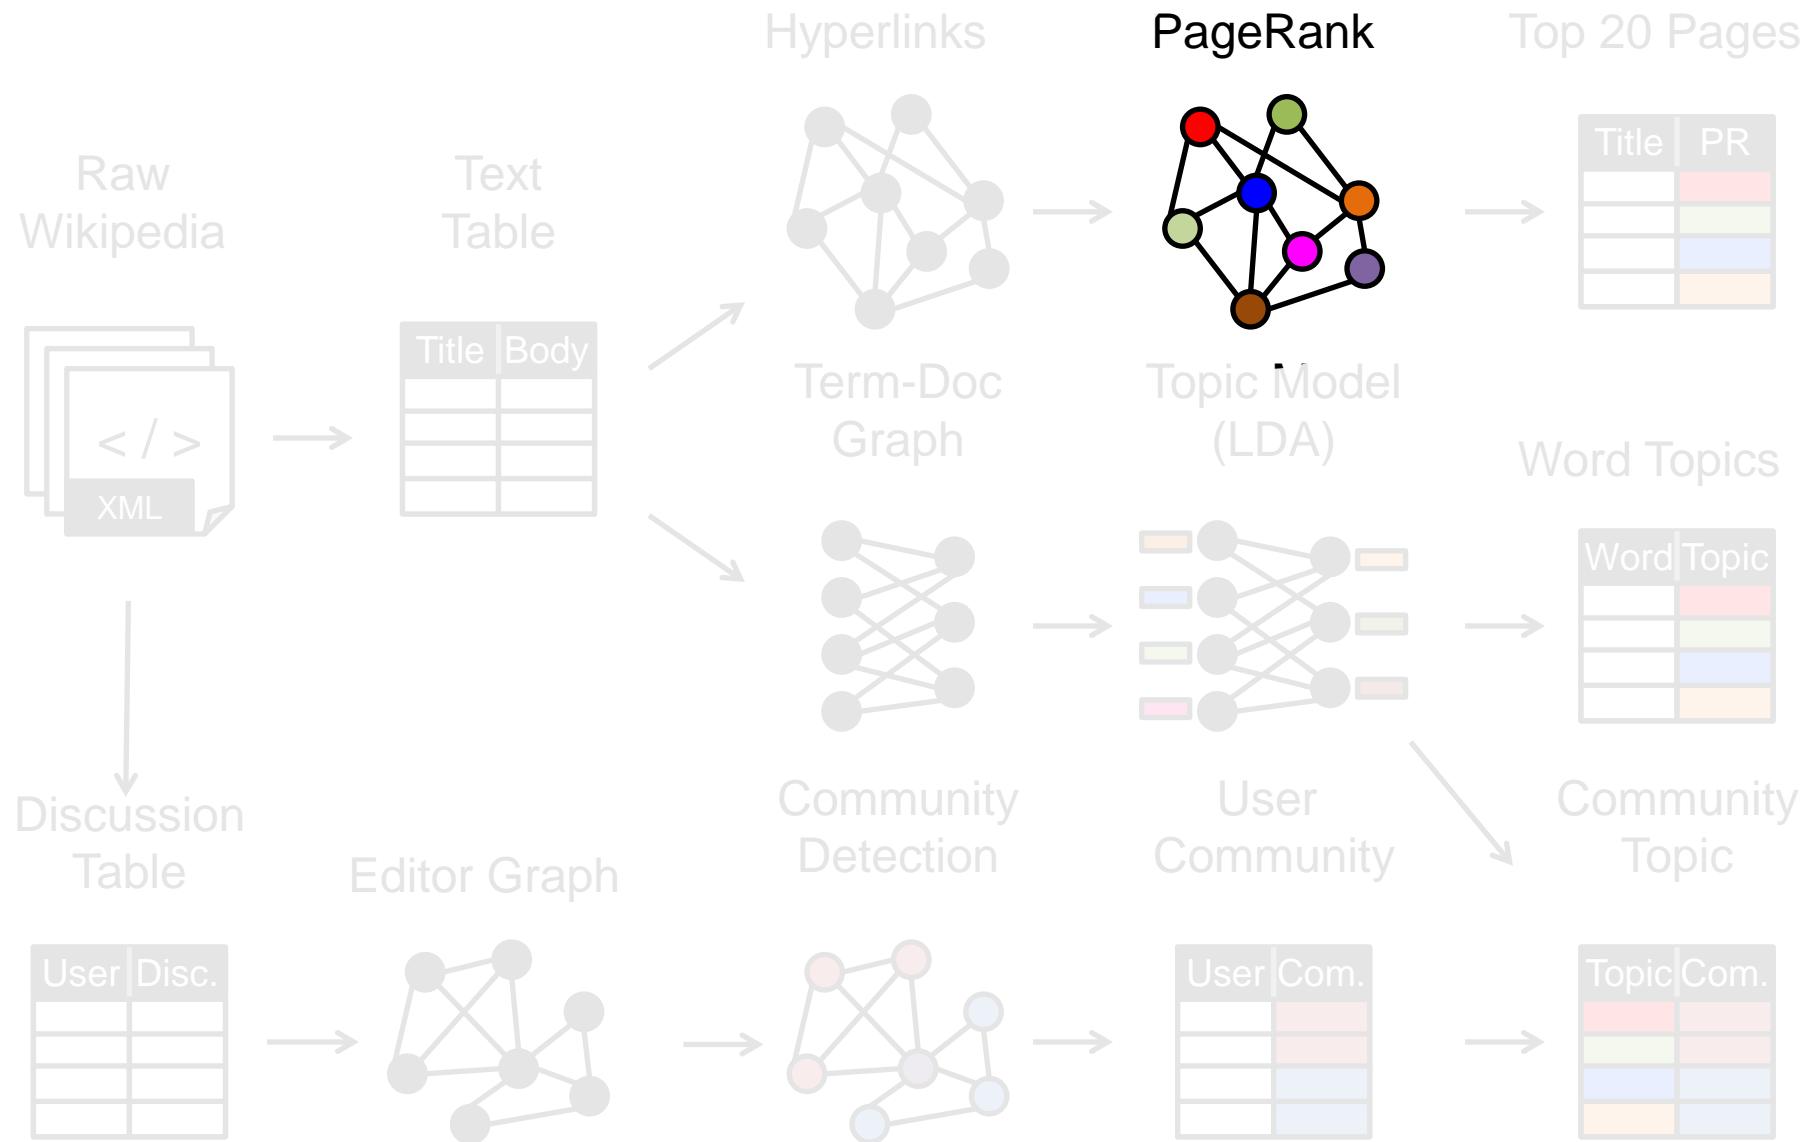
- PageRank
- Connected Components
- Triangle Counting Example



# Graphs are Central to Analytics

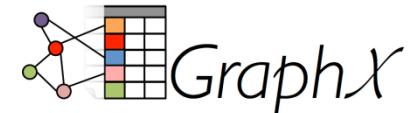


# Graphs are Central to Analytics



# Graphs are Essential to Data Mining and Machine Learning

- Identify influential people and information
- Find communities
- Understand people's shared interests
- Model complex data dependencies

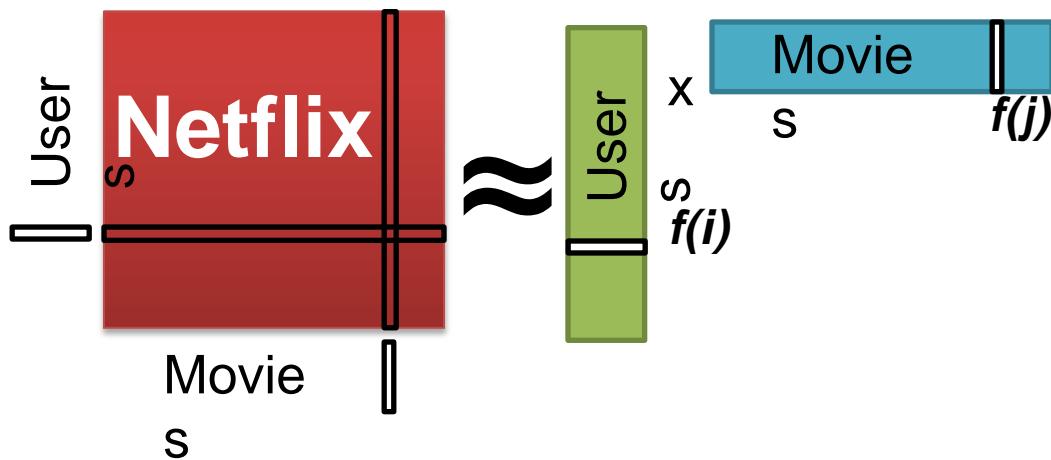


# Recommending Products



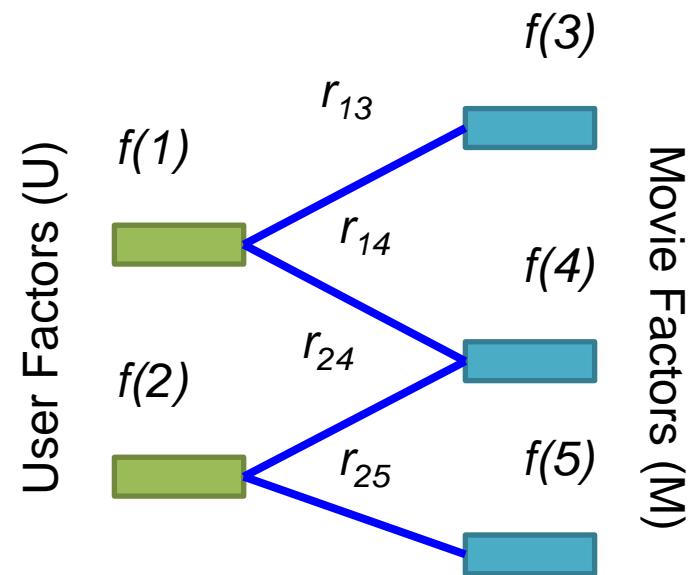
# Recommending Products

Low-Rank Matrix Factorization:



Iterate:

$$f[i] = \arg \min_{w \in \mathbb{R}^d} \sum_{j \in \text{Nbrs}(i)} (r_{ij} - w^T f[j])^2 + \lambda \|w\|_2^2$$

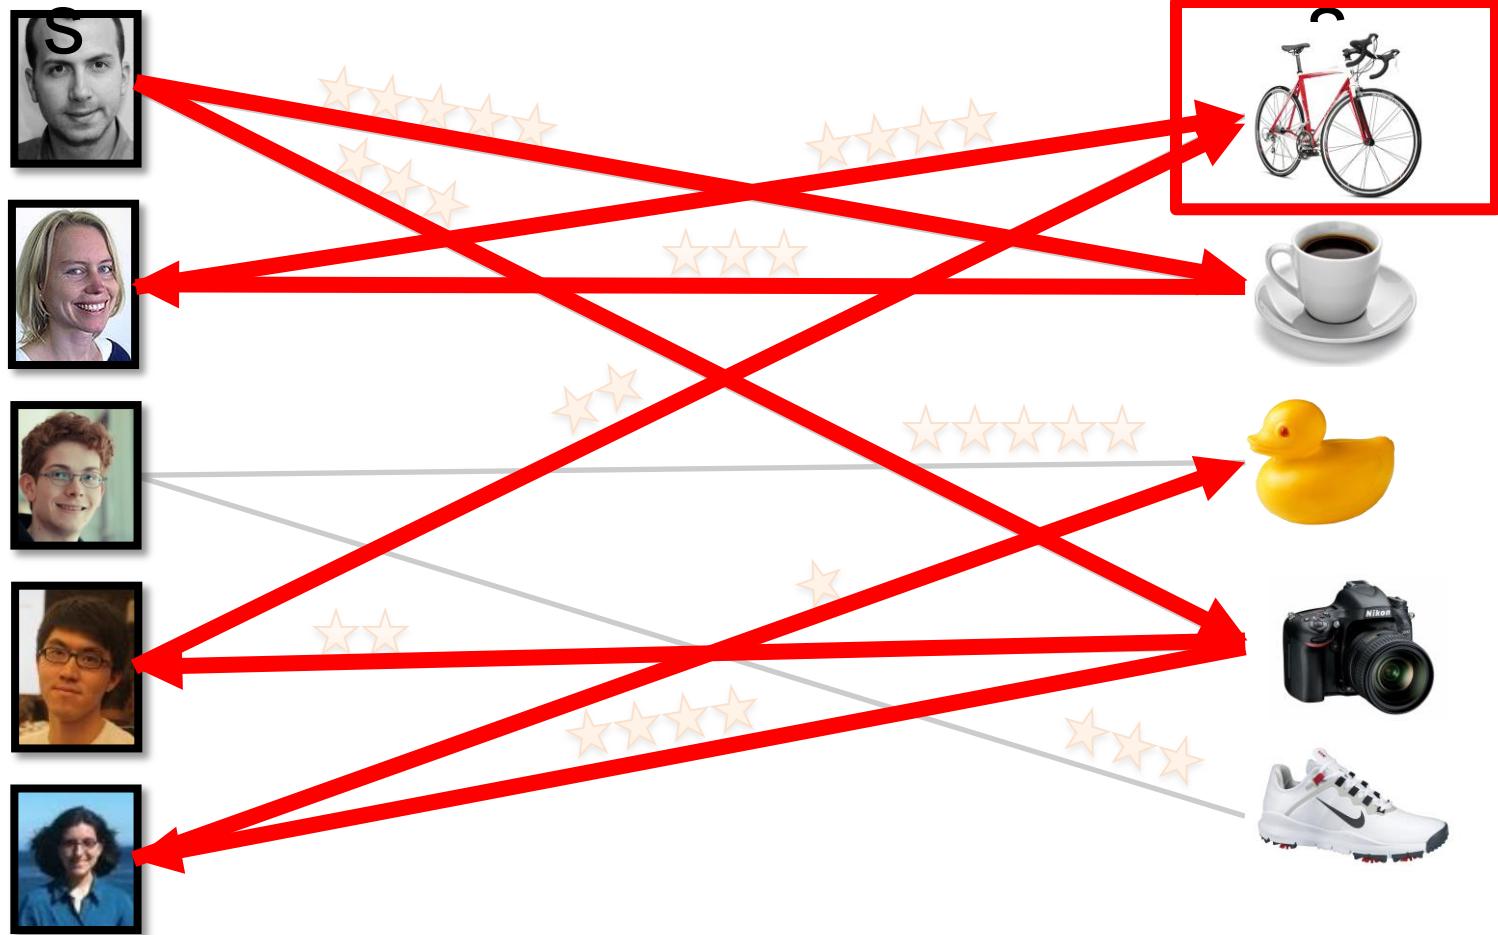


# Collaborative Filtering

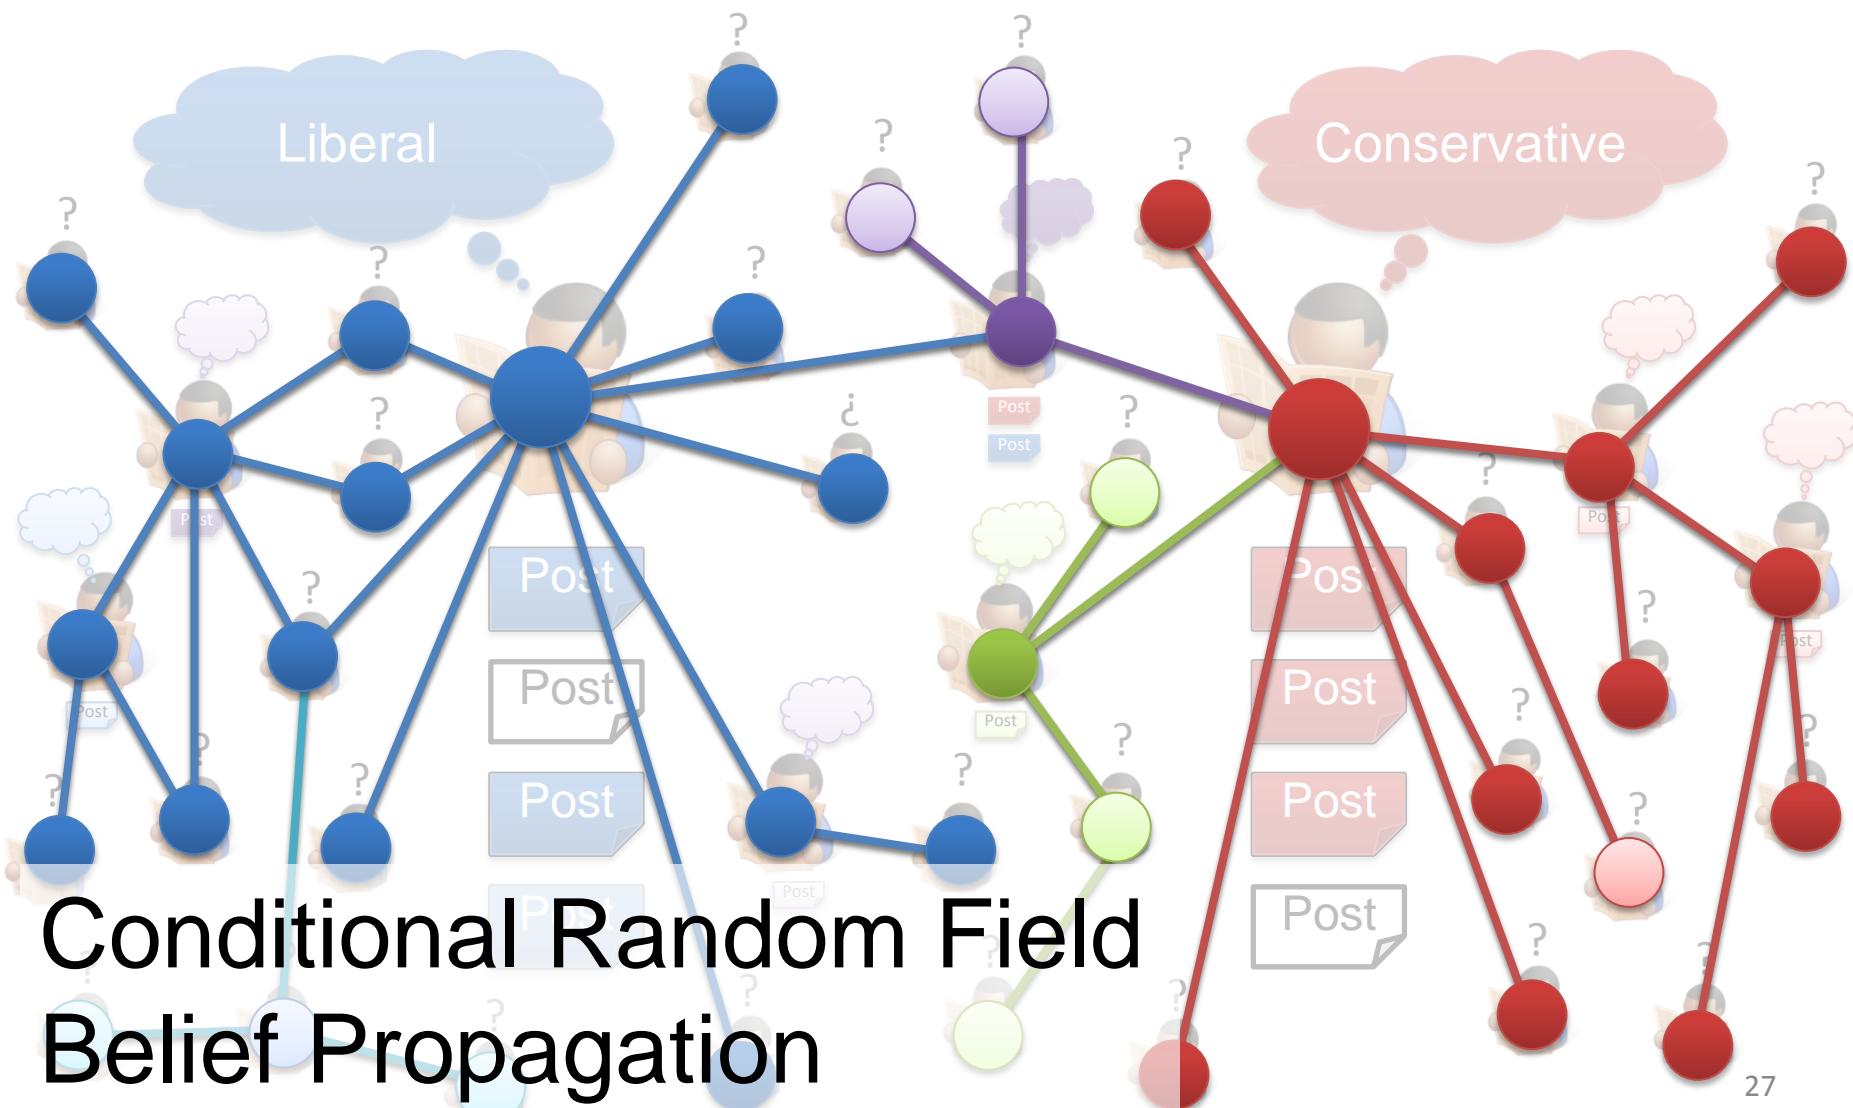
User

Ratings

Item

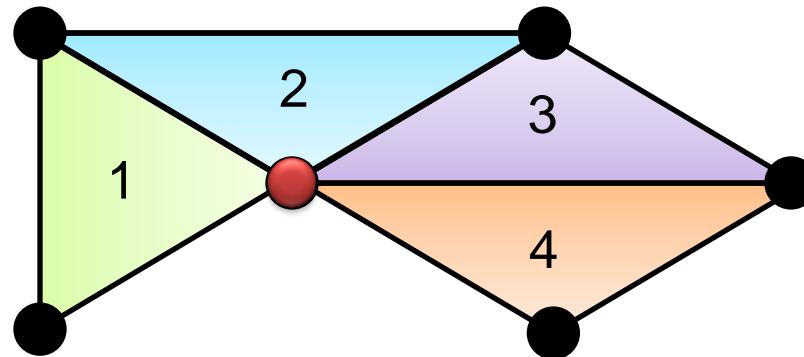


# *Predicting User Behavior*



# Finding Communities

- Count triangles passing through each vertex:



- Measures “cohesiveness” of local community



Fewer Triangles  
Weaker Community

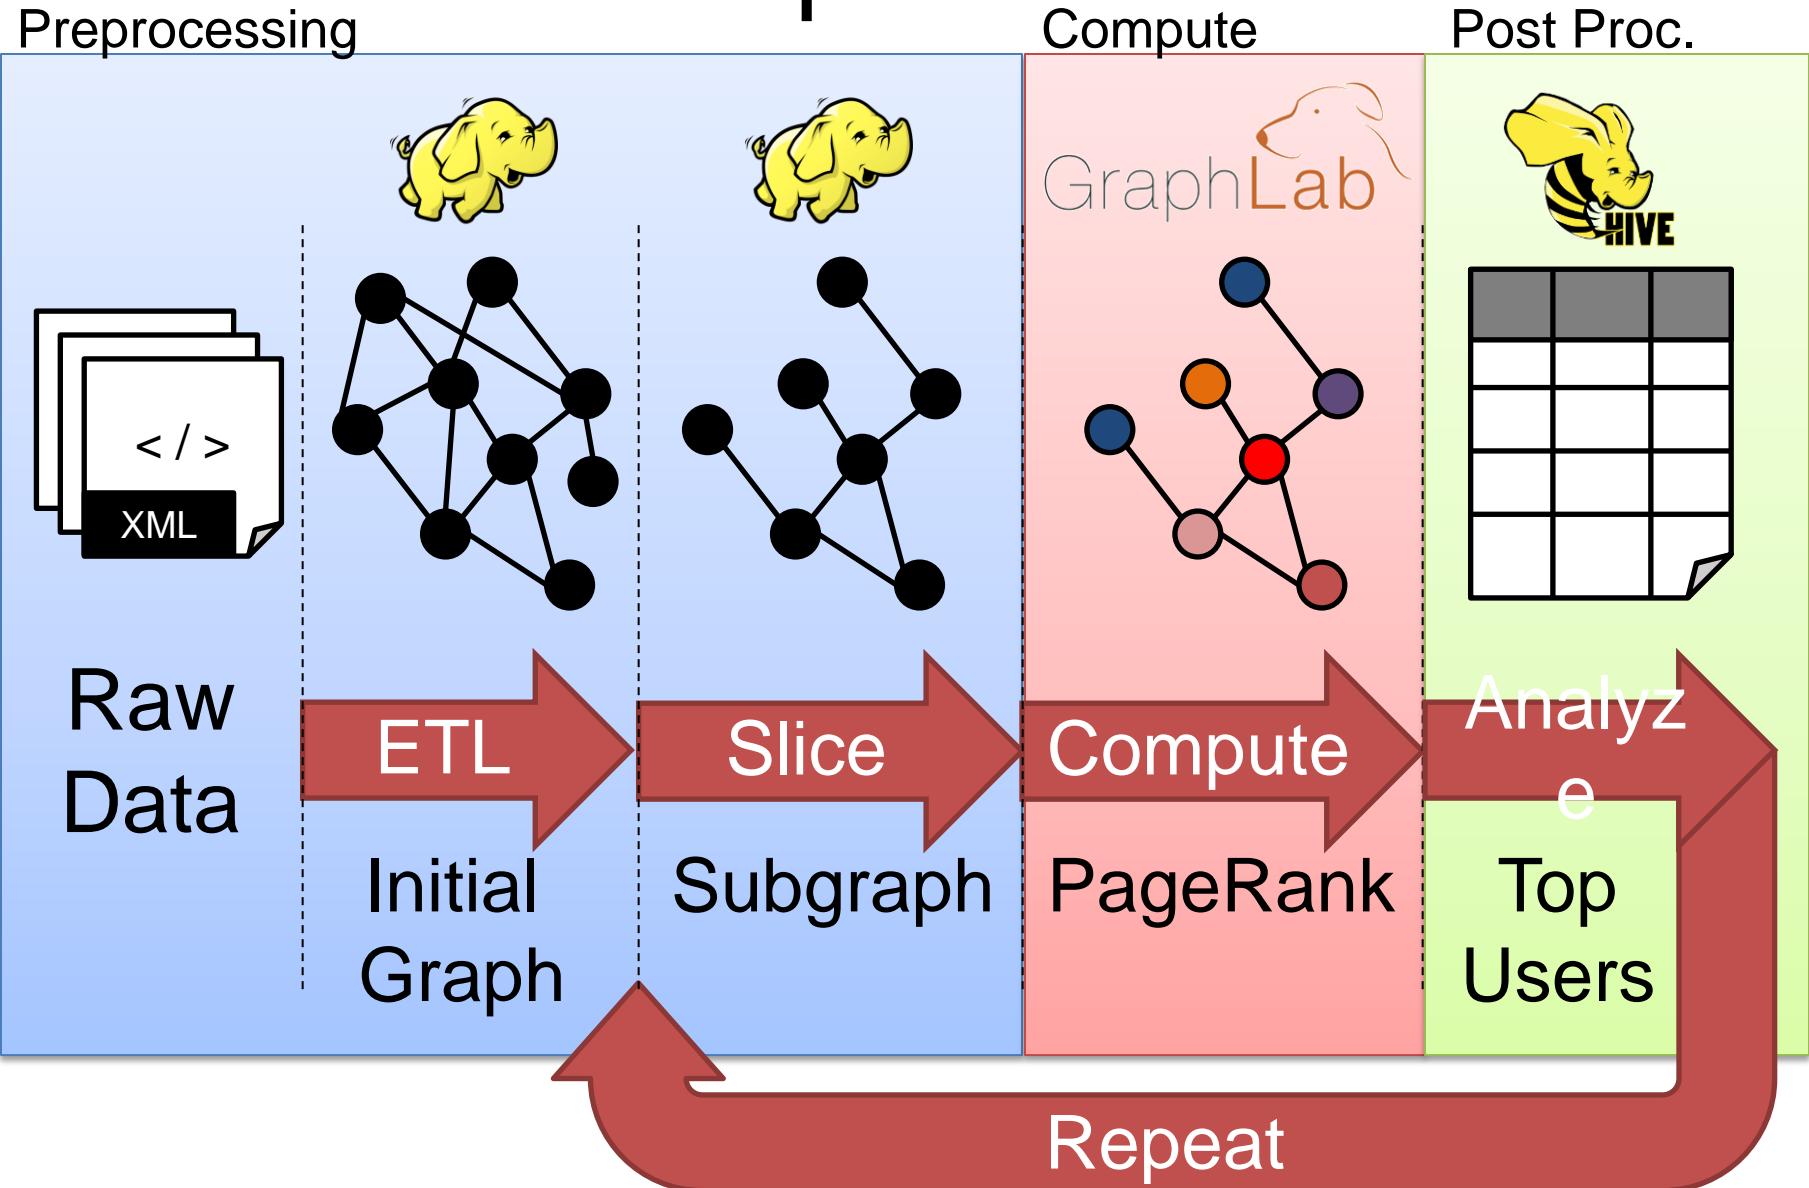


More Triangles  
Stronger Community

# Many More Graph Algorithms

- **Collaborative Filtering**
  - Alternating Least Squares
  - Stochastic Gradient Descent
  - Tensor Factorization
  - SVD
- **Structured Prediction**
  - Loopy Belief Propagation
  - Max-Product Linear Programs
  - Gibbs Sampling
- **Semi-supervised ML**
  - Graph SSL
- CoEM
- **Graph Analytics**
  - PageRank
  - Single Source Shortest Path
  - Triangle-Counting
  - Graph Coloring
  - K-core Decomposition
  - Personalized PageRank
- **Classification**
  - Neural Networks
  - Lasso
- ...

# Example Graph Analytics Pipeline



# Separate Systems to Support Each View

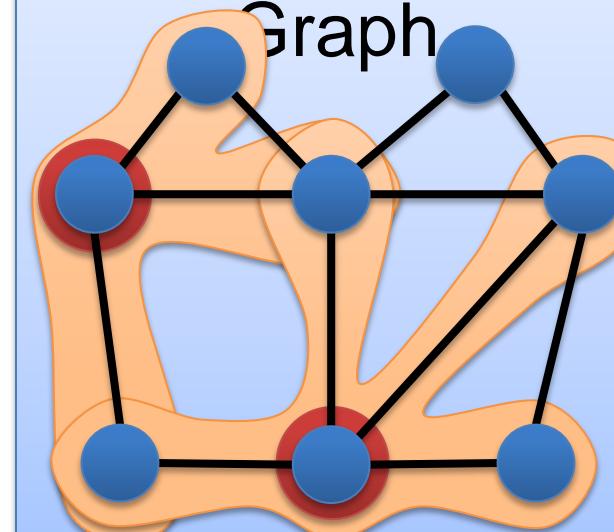
## Table View



## Graph View



### Dependency Graph



## GraphX – Adding Graphs to the Mix

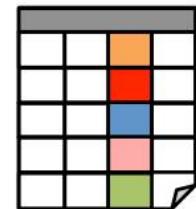
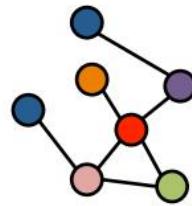
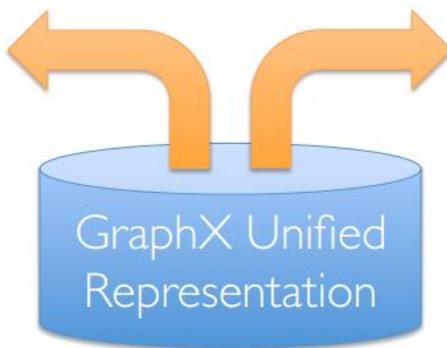


Table View



Graph View

Tables and Graphs are *composable views* of the same *physical data*

Each view has its own *operators* that *exploit the semantics* of the view to achieve *efficient execution*

R. Xin, J. Gonzalez, M. Franklin, I. Stoica, "GraphX: In-situ Graph Computation Made Easy"  
GRADES Workshop at SIGMOD, June 2013.

## GraphX: Graphs → Dataflow

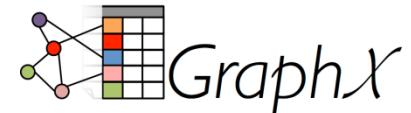
1. Encode graphs as distributed tables
2. Express graph computation in relational ops.
3. Recast graph systems optimizations as:
  1. Distributed join optimization
  2. Incremental materialized maintenance

---

Integrate Graph and  
Table data processing  
systems.

Achieve performance  
parity with specialized  
systems.

*Having separate systems  
for each view is  
**difficult to use and inefficient***



# Difficult to Program and Use

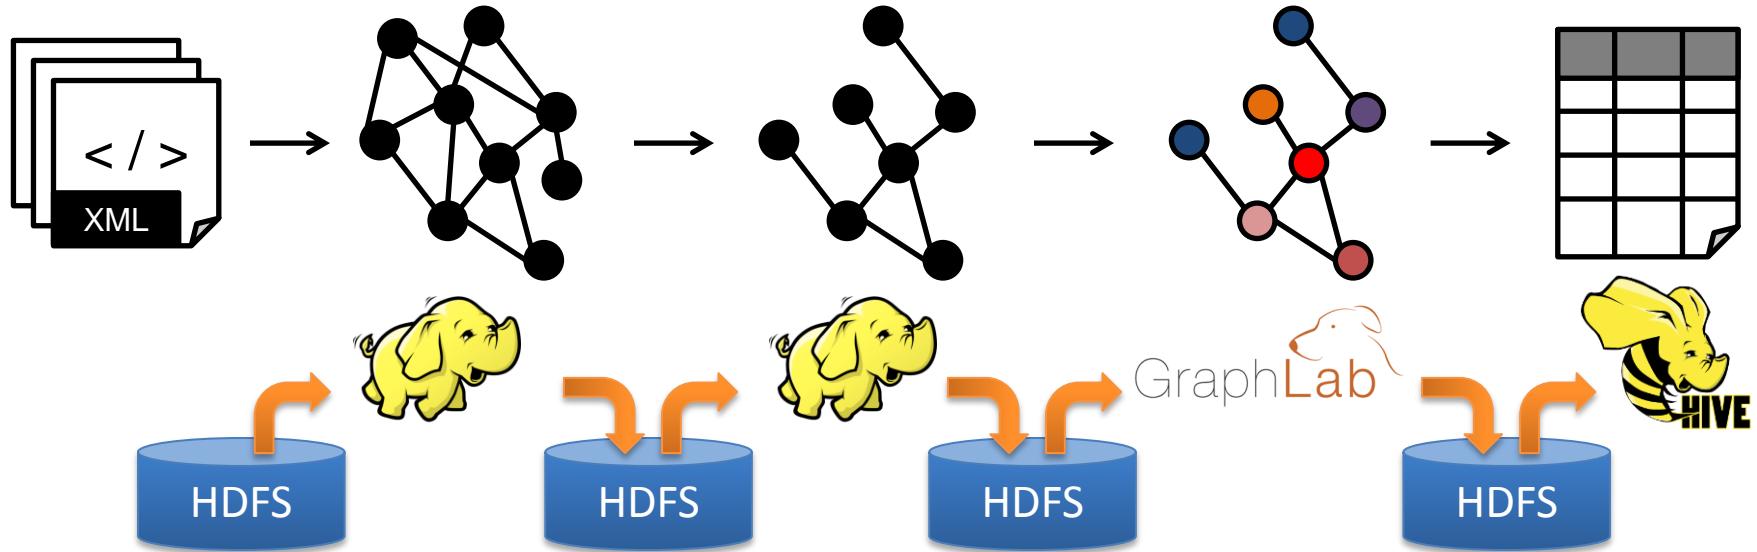
Users must *Learn*, *Deploy*, and  
*Manage* multiple systems



Leads to brittle and often  
complex interfaces

# Inefficient

Extensive **data movement** and **duplication** across the network and file system

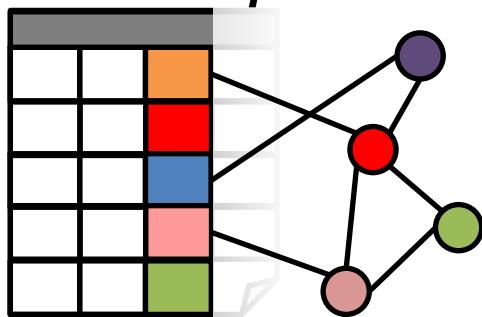


Limited reuse internal data-structures across stages

# *Solution:* The GraphX Unified Approach

## New API

*Blurs the distinction  
between Tables and  
Graphs*



## New System

*Combines Data-Parallel  
Graph-Parallel Systems*



Enabling users to **easily** and **efficiently**  
express the entire graph analytics  
pipeline

Tables and Graphs are **composable views** of the *same physical data*

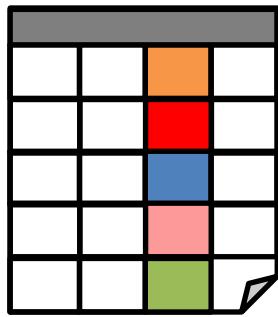
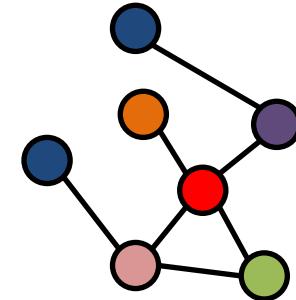
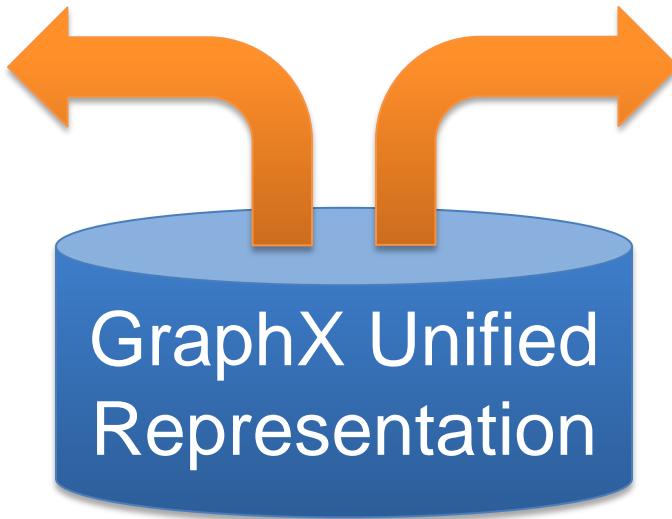


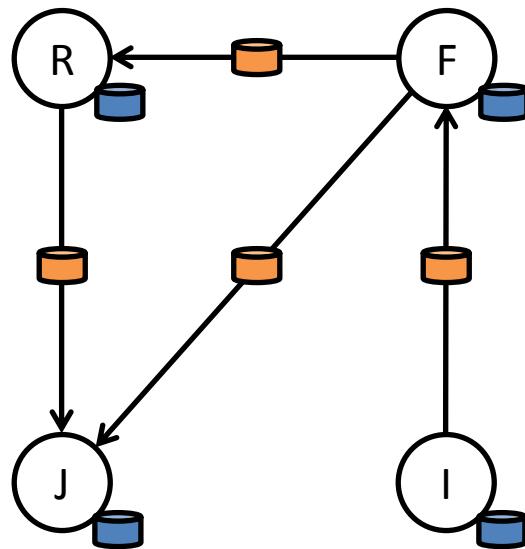
Table View



Graph View

Each view has its own **operators** that  
**exploit the semantics** of the view  
to achieve **efficient execution**

# Property Graph



# View a Graph as a Table

## Vertex Property Table

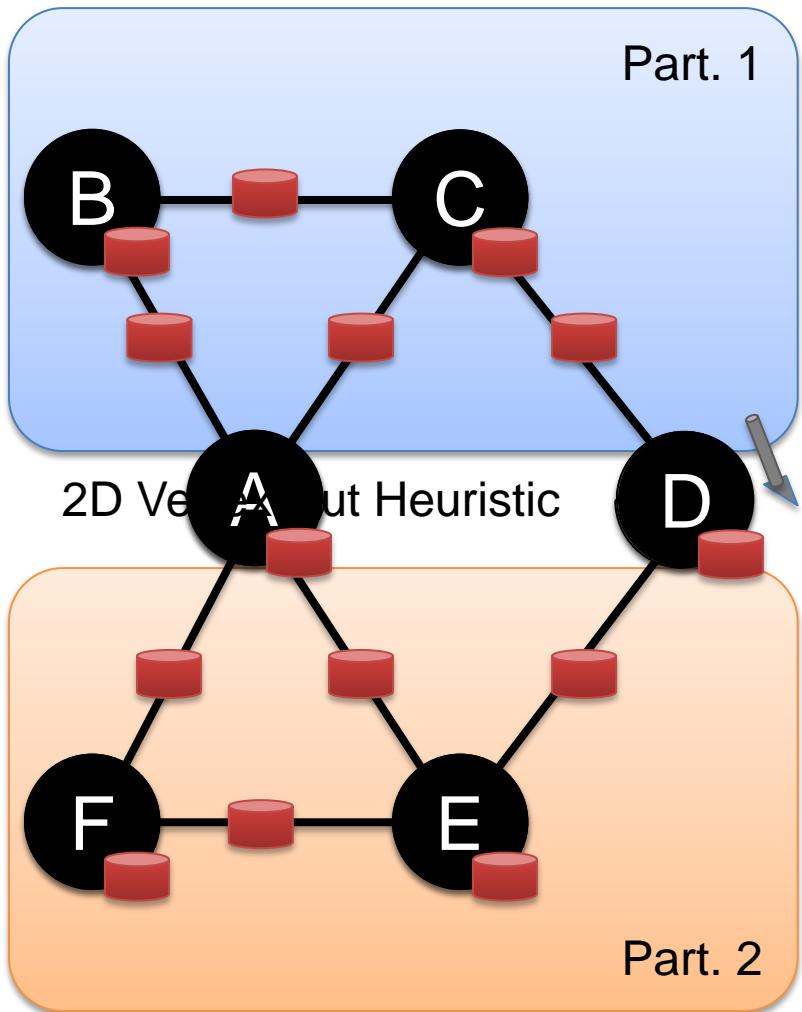
Id	Property (V)
Rxin	(Stu., Berk.)
Jegonzal	(PstDoc, Berk.)
Franklin	(Prof., Berk)
Istoica	(Prof., Berk)

## Edge Property Table

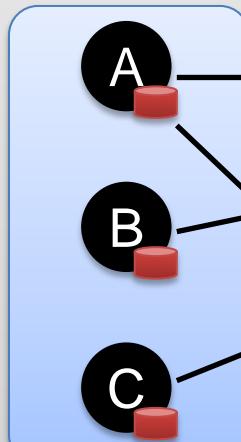
SrcId	DstId	Property (E)
rxin	jegonzal	Friend
franklin	rxin	Advisor
istoica	franklin	Coworker
franklin	jegonzal	PI

# Distributed Graphs as Tables (RDDs)

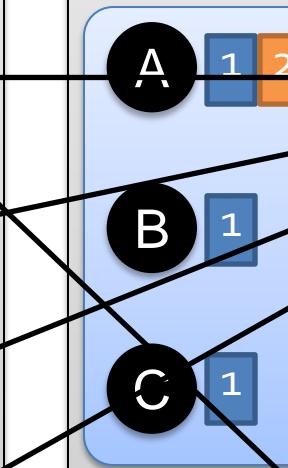
Property Graph



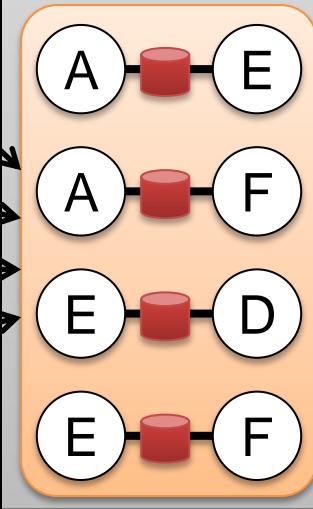
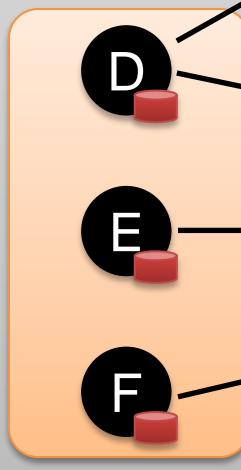
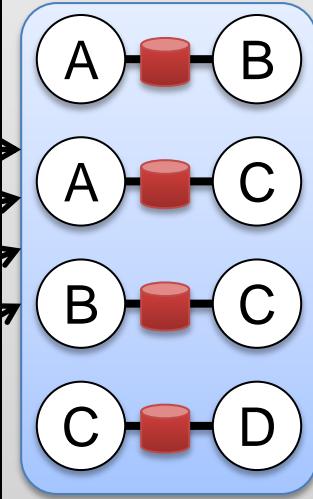
Vertex Table (RDD)



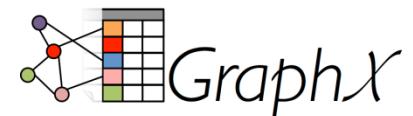
Routing Table (RDD)



Edge Table (RDD)



*By exploiting **graph-structure**  
Graph-Parallel systems  
can be **orders-of-magnitude**  
**faster.***



# Triangle Counting on Twitter

40M Users, 1.4 Billion Links

Counted: 34.8 Billion  
Triangles

Hadoop  
[WWW'11]

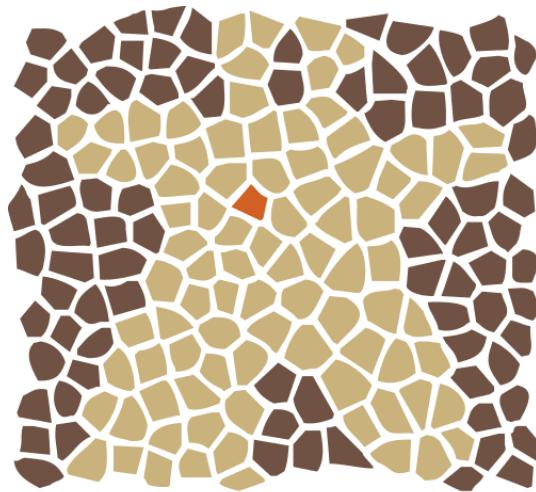
1536 Machines  
423 Minutes

GraphLab

64 Machines  
15 Seconds

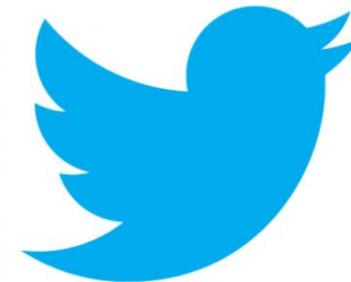
1000 x  
Faster

# Specialized Graph Systems



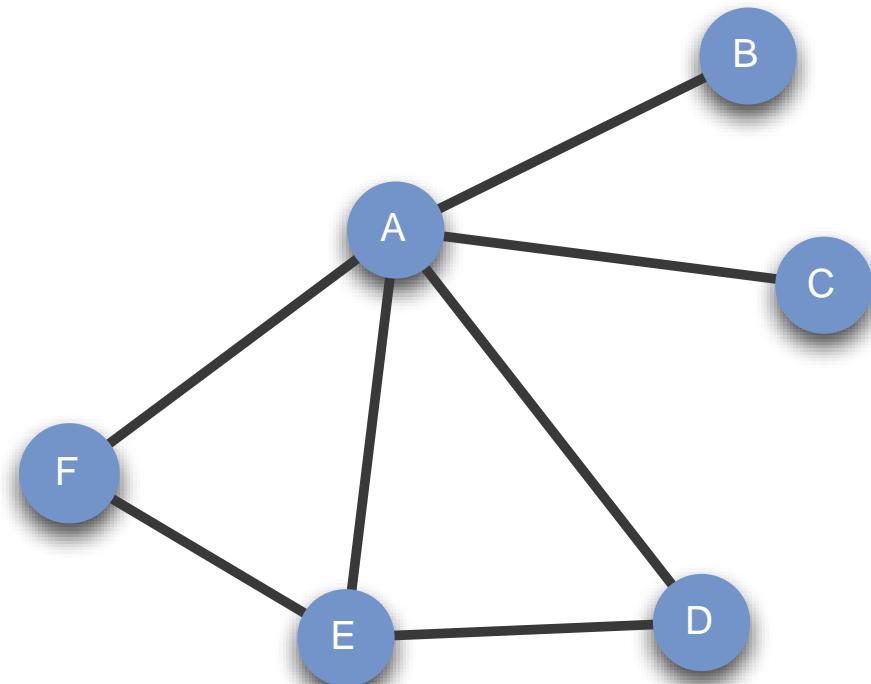
APACHE  
GIRAPH

Google Pregel



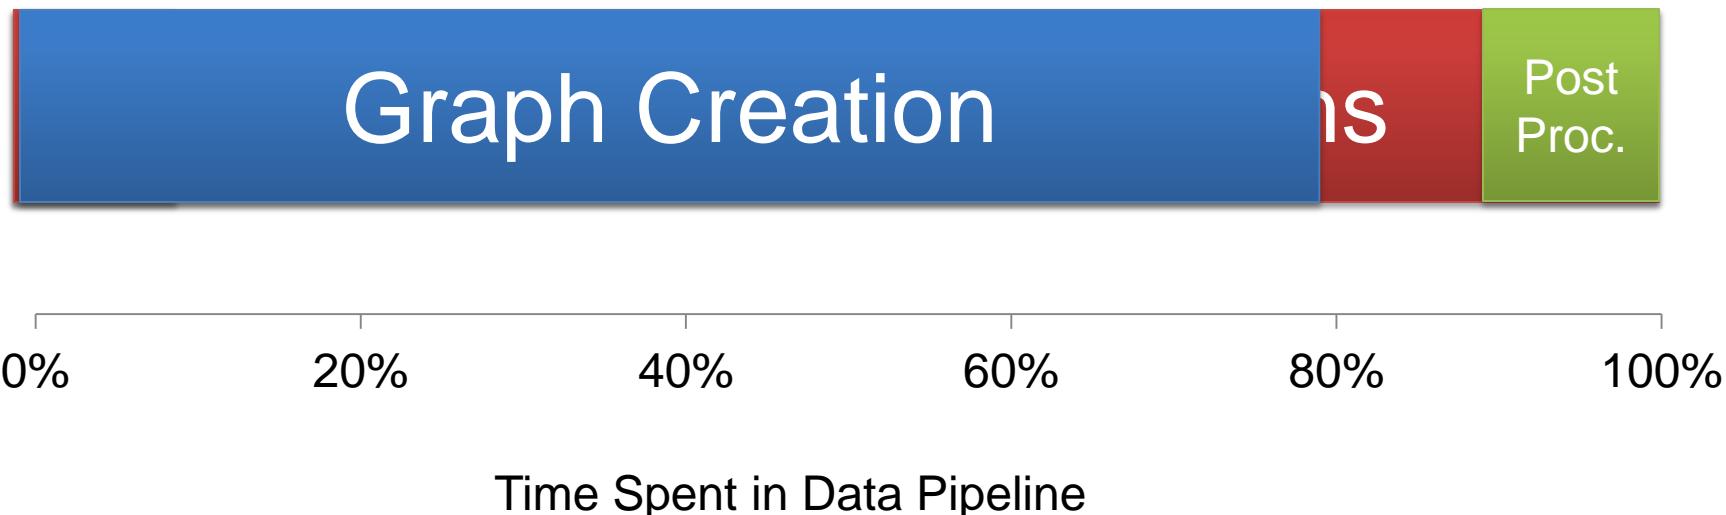
# Specialized Graph Systems

1. APIs to capture complex graph dependencies
2. Exploit graph structure to reduce communication and computation



# *Why GraphX?*

# The Bigger Picture



# GraphX Unifies Data-Parallel and Graph-Parallel Systems

Spark

*Table API*

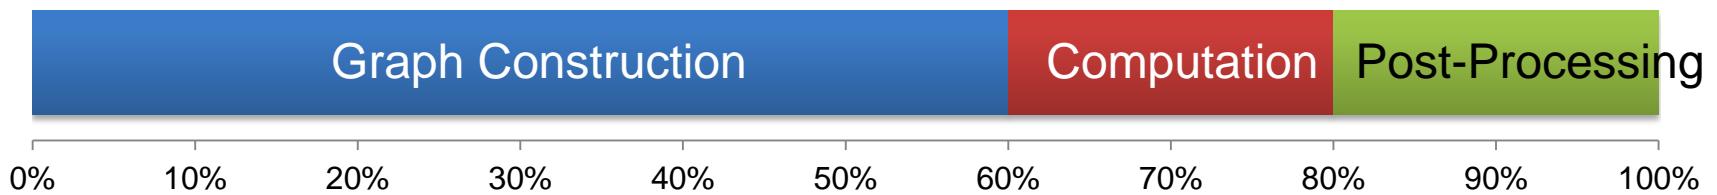
RDDs, Fault-tolerance, and task scheduling

one system for the entire graph pipeline

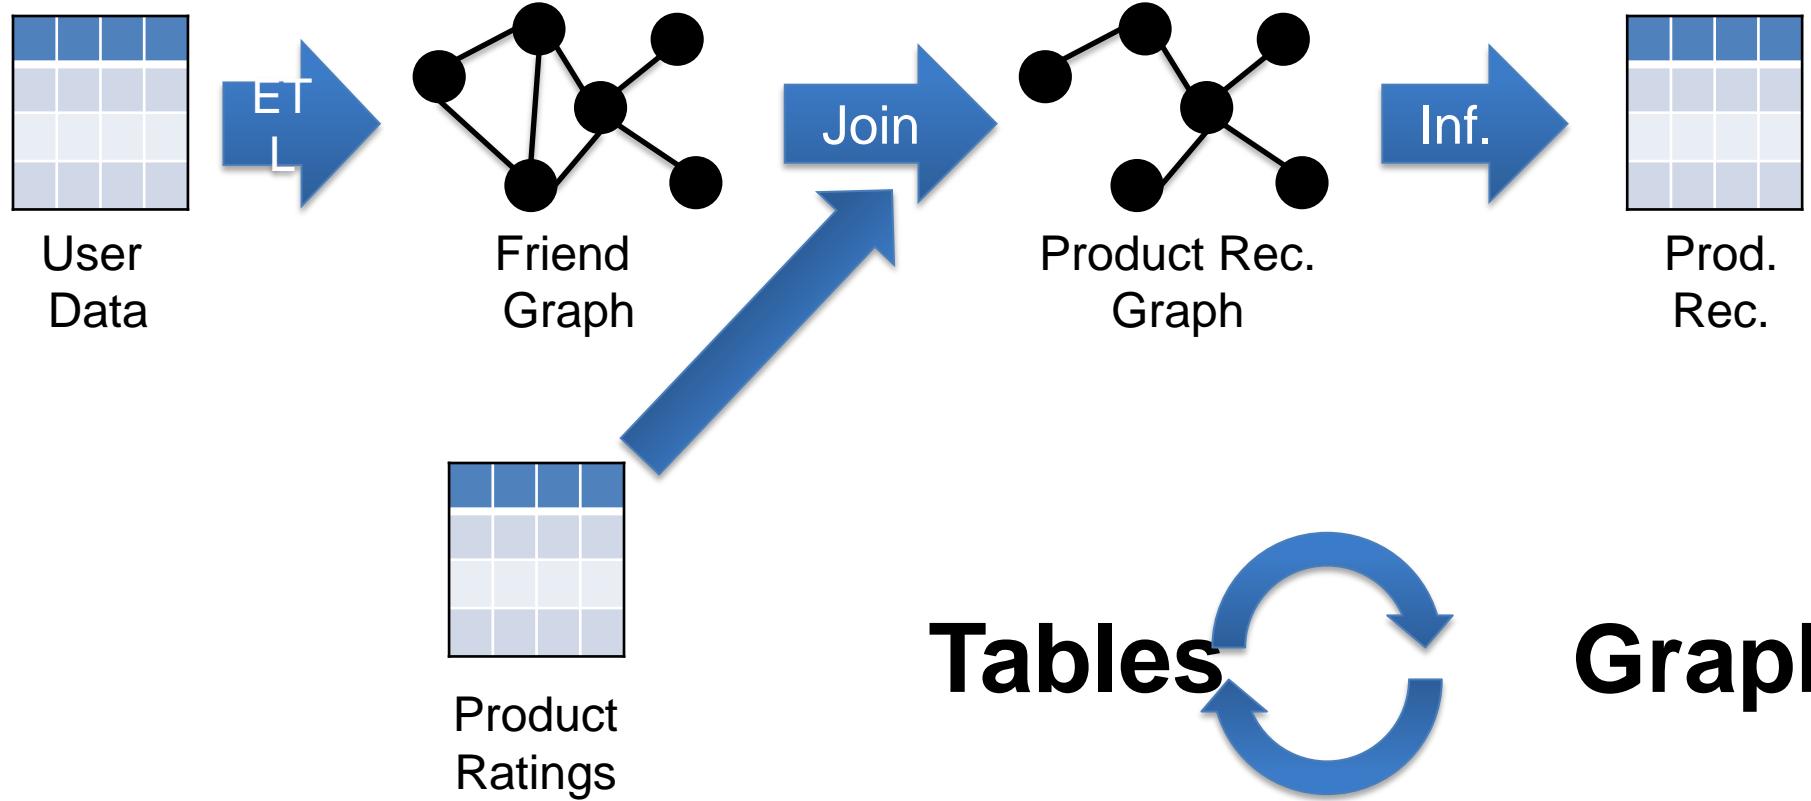
GraphLab

*Graph API*

graph representation and execution



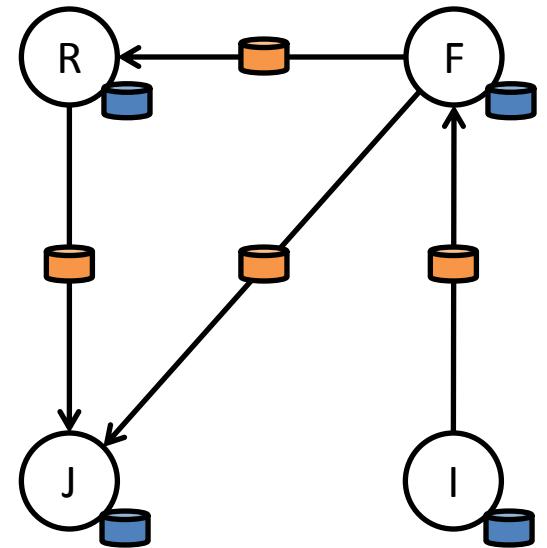
# Enable Joining Tables and Graphs



# The GraphX Resilient Distributed Graph

Id	Attribute (V)
Rxin	(Stu., Berk.)
Jegonzal	(PstDoc, Berk.)
Franklin	(Prof., Berk)
Istoica	(Prof., Berk)

SrcId	DstId	Attribute (E)
rxin	jegonzal	Friend
franklin	rxin	Advisor
istoica	franklin	Coworker
franklin	jegonzal	PI



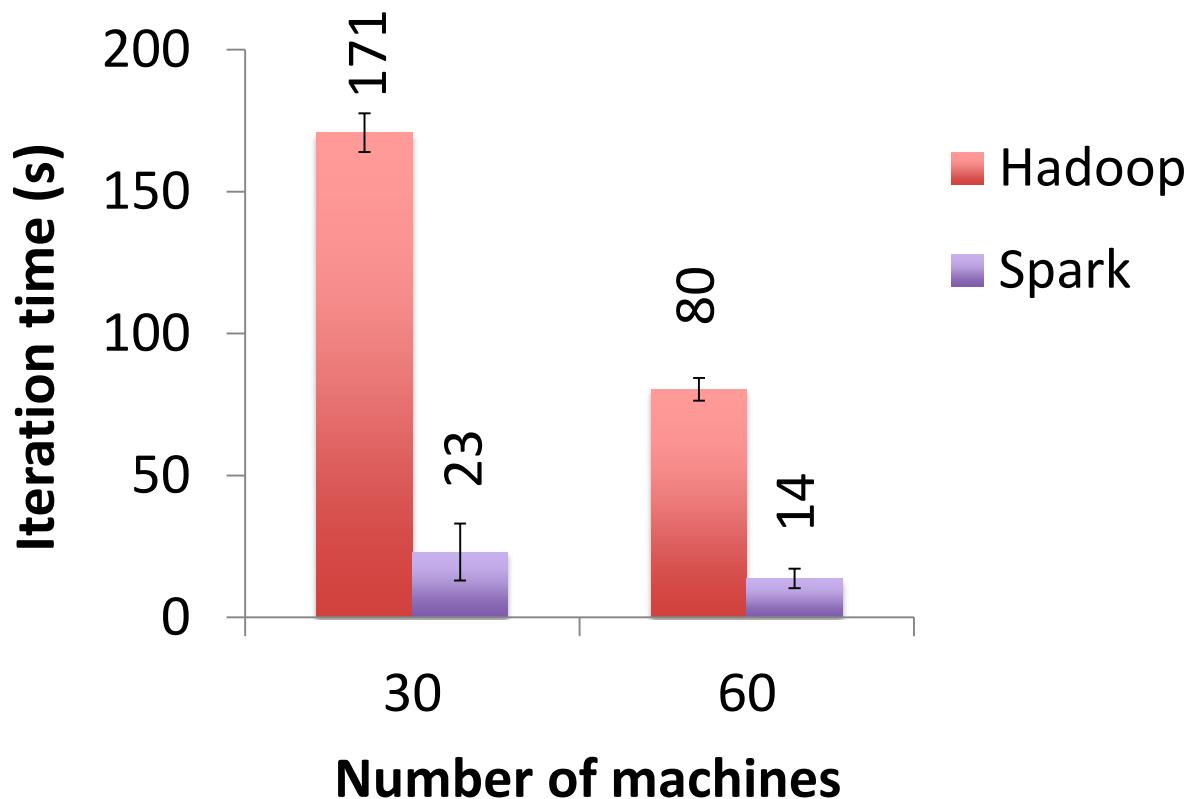
# Table Operators

- Table (RDD) operators are inherited from Spark:

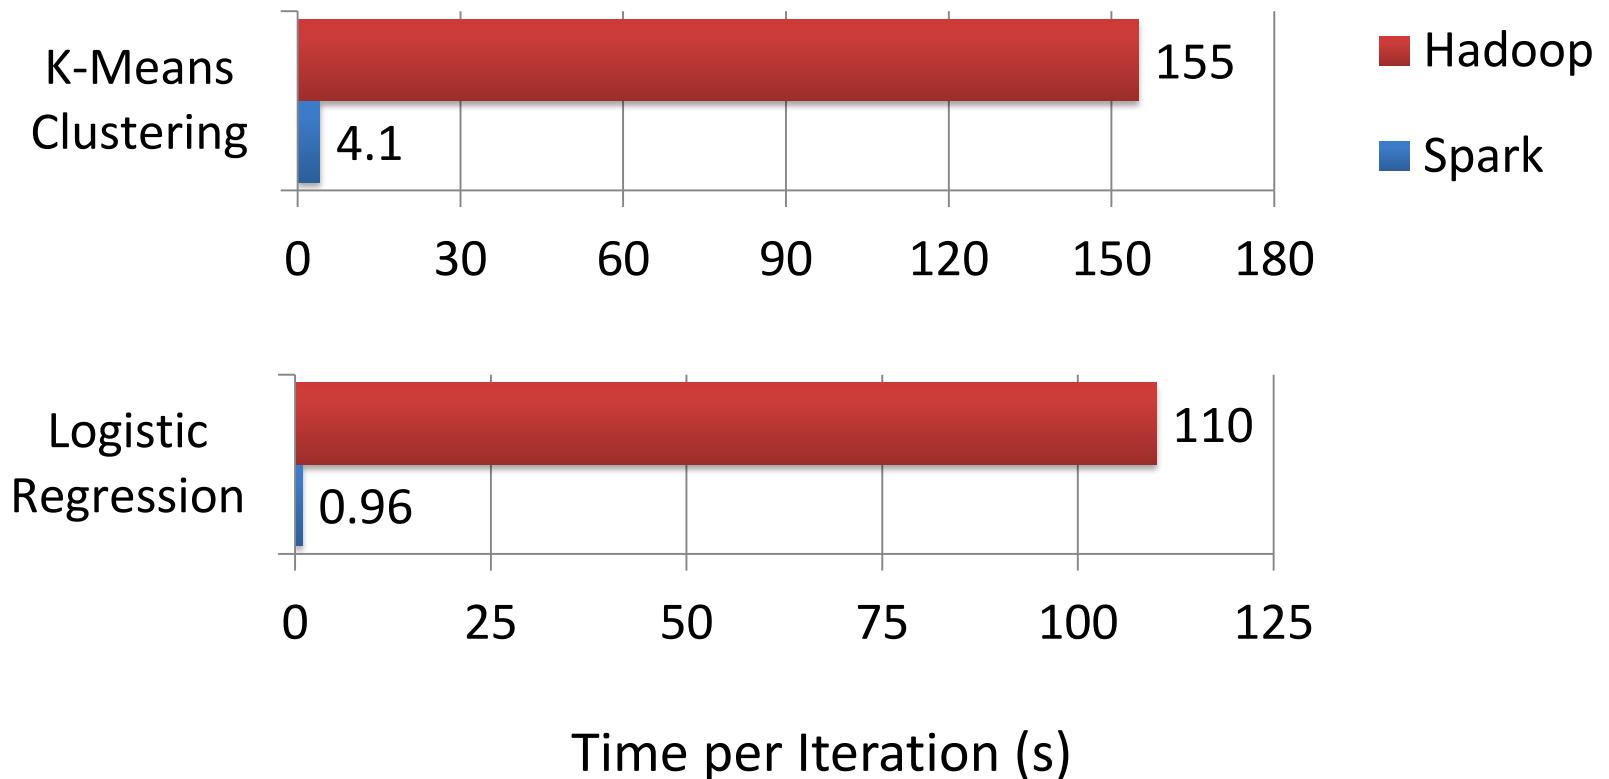
map	reduce	sample
filter	count	take
groupBy	fold	first
sort	reduceByKey	partitionBy
union	groupByKey	mapwith
join	cogroup	pipe
leftOuterJoin	cross	save
rightOuterJoin	zip	...

# -GraphX API

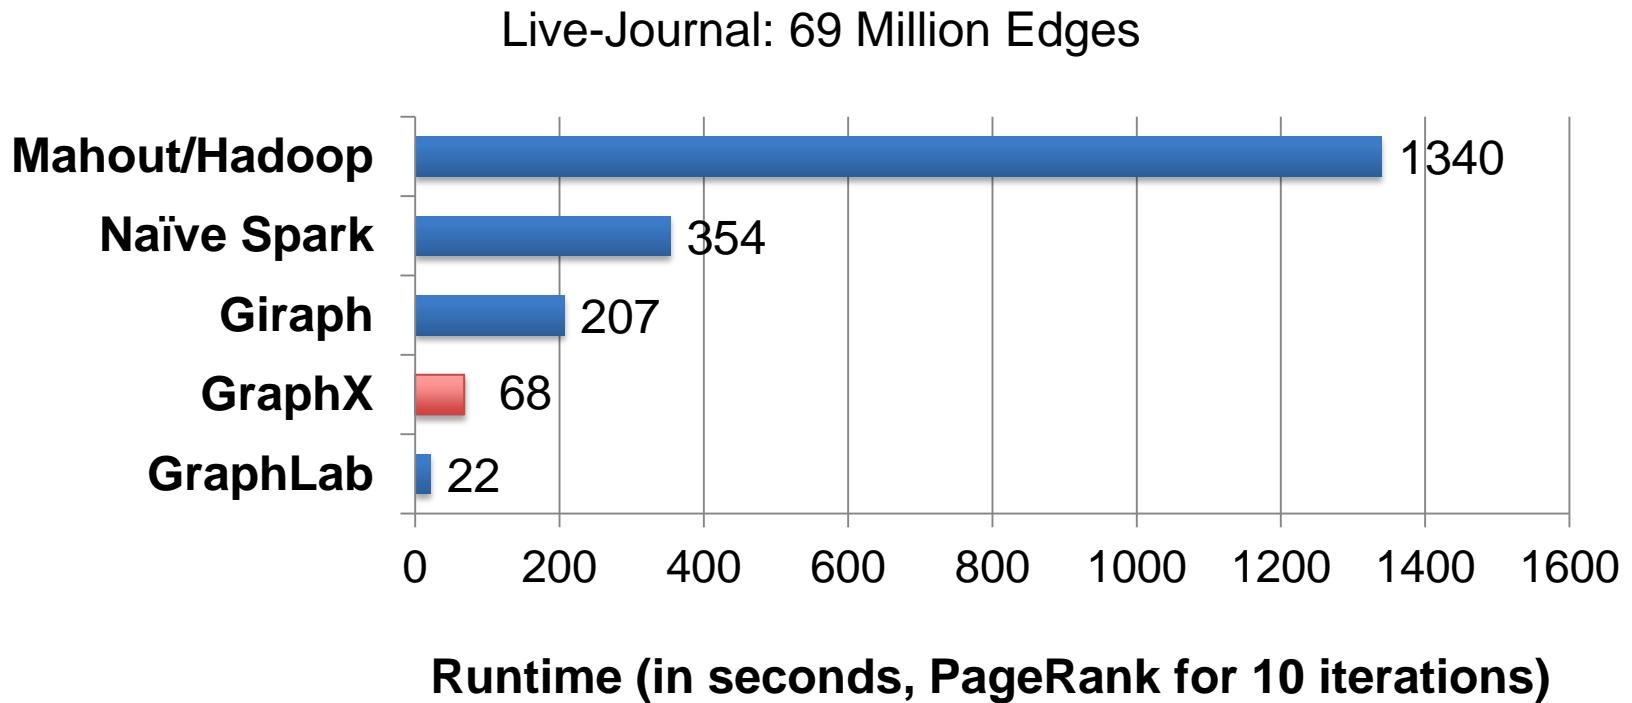
# PageRank Performance



# Other Iterative Algorithms



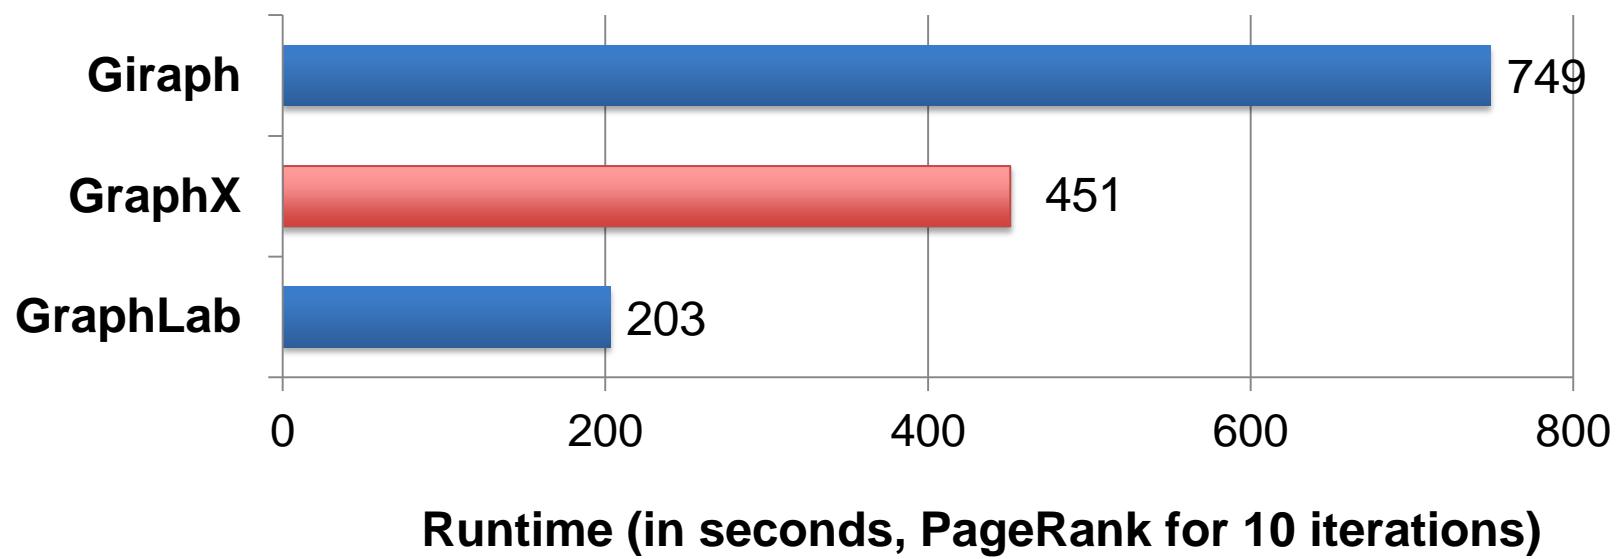
# Performance Comparisons



GraphX is roughly **3x slower** than GraphLab

# GraphX scales to larger graphs

Twitter Graph: 1.5 Billion Edges



GraphX is roughly 2x slower than GraphLab

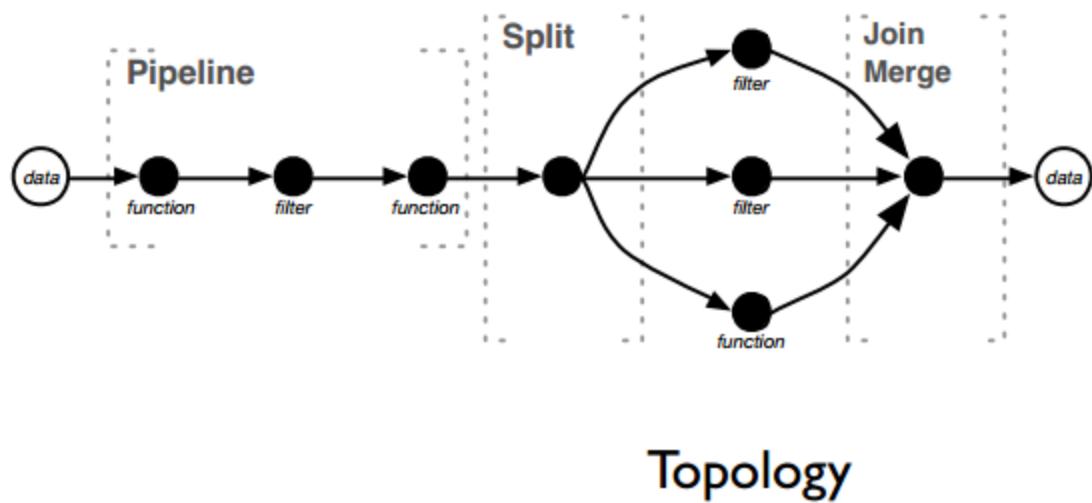
- » Scala + Java overhead: Lambdas, GC time, ...
- » No shared memory parallelism: 2x increase in comm.

PageRank is just one stage....

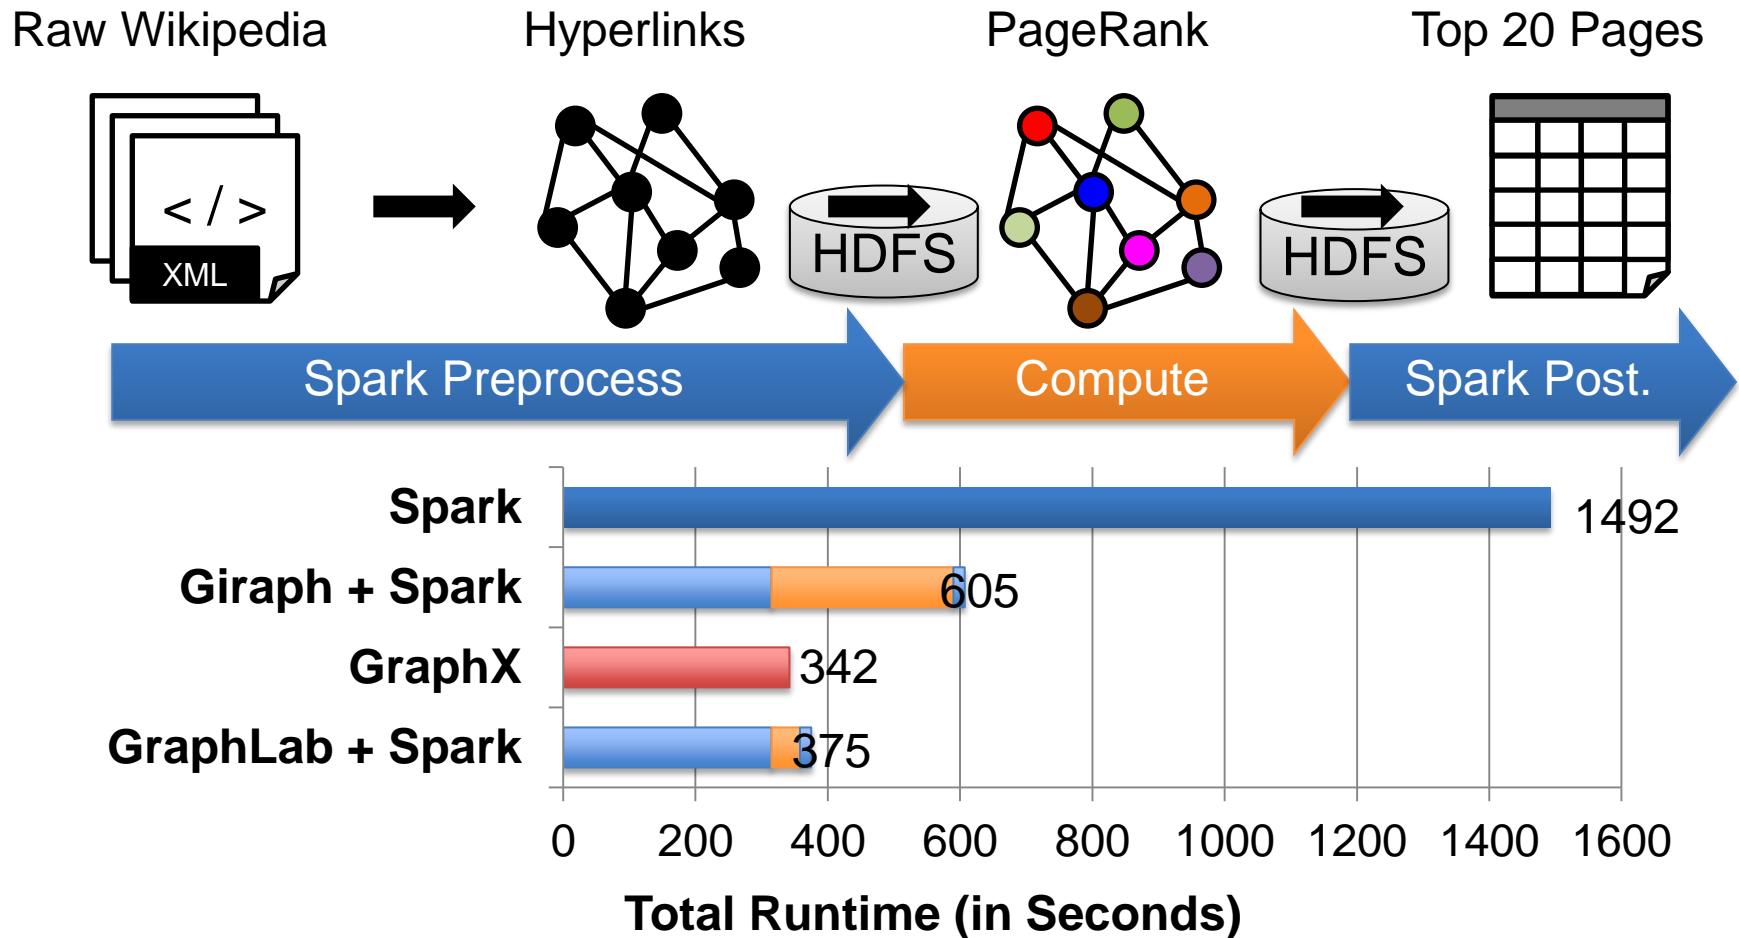
What about a [pipeline](#)?

# SOME COMMON PATTERNS

- Functions
- Filters
- Joins
  - Inner / Outer / Mixed
  - Asymmetrical / Symmetrical
- Merge (Union)
- Grouping
  - Secondary Sorting
  - Unique (Distinct)
- Aggregations
  - Count, Average, etc



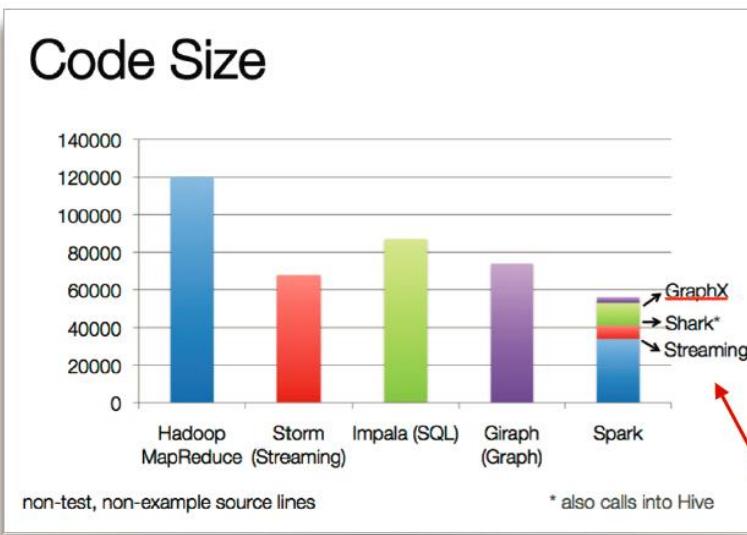
# A Small Pipeline in GraphX



Timed end-to-end GraphX is *faster* than  
GraphLab

# Spark v Hadoop - Code Size

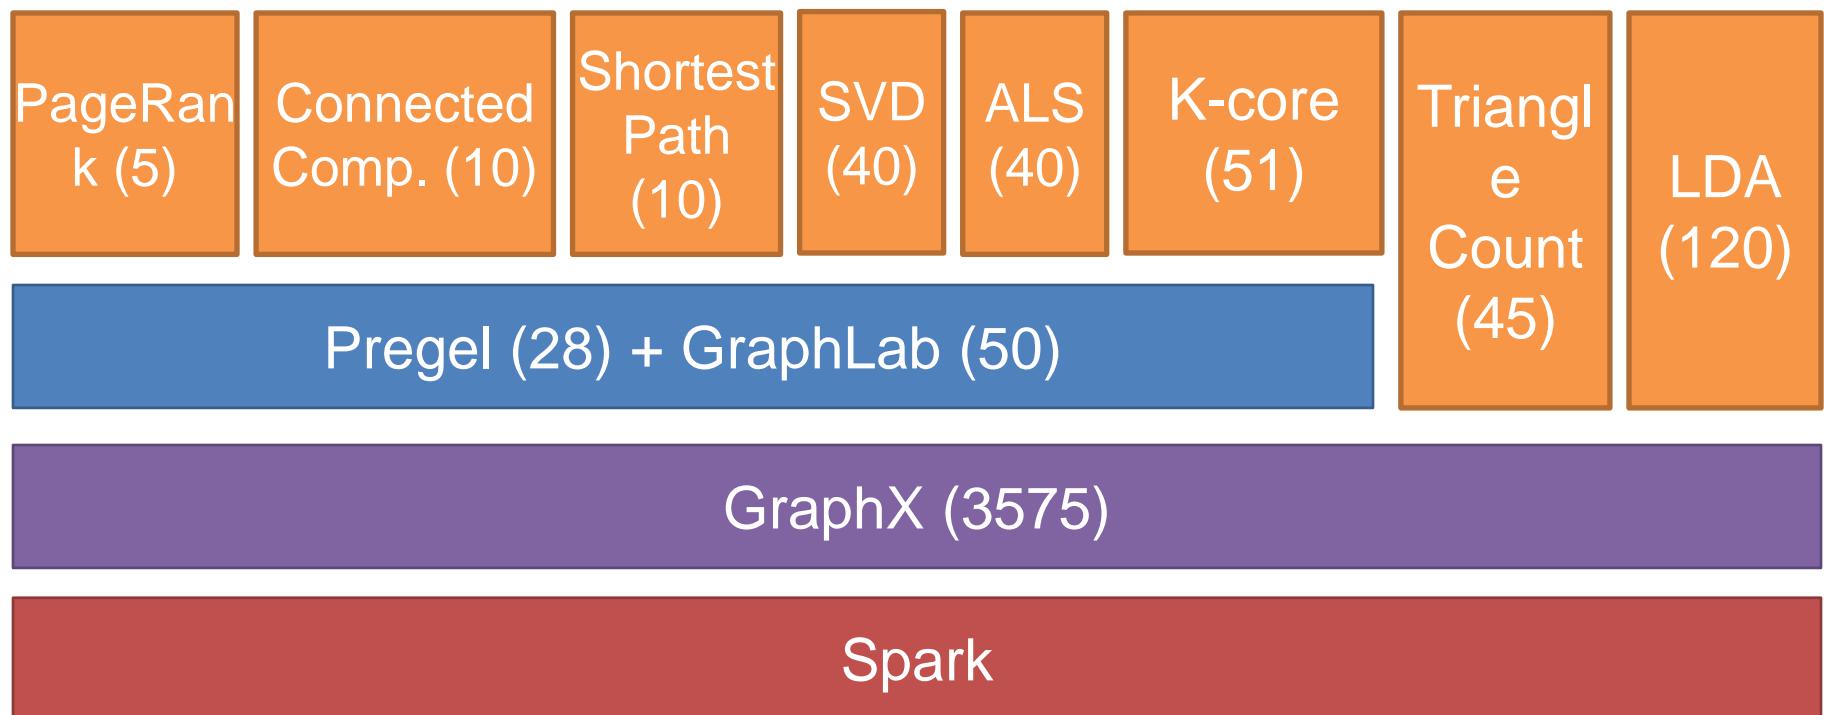
## A Brief History: Spark



*The State of Spark, and Where We're Going Next*  
Matei Zaharia  
Spark Summit (2013)  
[youtu.be/nU6vO2EJAb4](https://youtu.be/nU6vO2EJAb4)

*used as libs, instead of  
specialized systems*

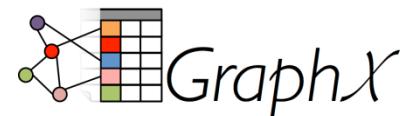
# The GraphX Stack (Lines of Code)



# Graph Analytics Code Example

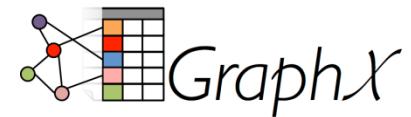
# Code Example – Use Case

- Use Case:
  - Use GraphX to analyze Wikipedia data and implement graph algorithms in Spark.
- Data Sets:
  - Articles (Nodes/Vertices): `graphx-wiki-vertices.txt`
  - Links to articles (Edges/Relationships): `graphx-wiki-edges.txt`



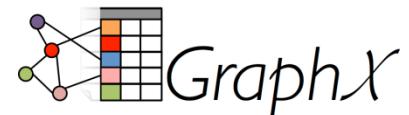
# Code Example Steps

- Import Spark GraphX API
- Load Article data from text file into RDD
- Load Links data from text file into RDD
- Construct the Graph using the two datasets
- Parse article rows into pairs of vertex ID and title
- Parse link rows into Edge objects with placeholder 0 attribute
- Create graph by calling Graph constructor with vertex RDD, our edge RDD & default vertex attribute
- Cache the resulting graph in memory to avoid reloading it from disk each time we use it
- Verify the data created so far



# Code Example Steps - 2

- Run graph analytics\*\*
- Run PageRank on Wikipedia
- Find titles of ten most important articles in Berkeley subgraph of Wikipedia



# Demo

# Serialization

- GraphX requires Kryo serializer to achieve maximum performance
- spark.serializer property in Spark UI
- By default Kryo Serialization is not enabled
- Enable Kryo Serialization for the spark shell
- Set Serialization setting in "spark-shell.cmd" file

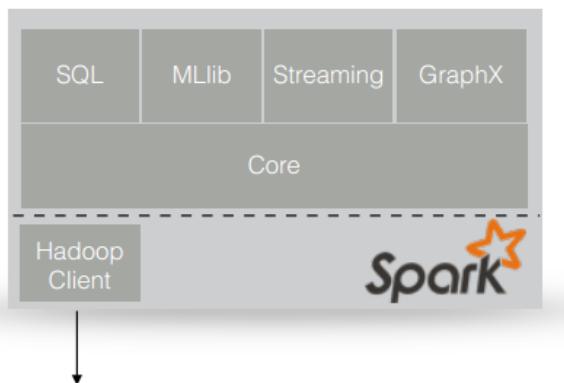


# Demo with Serialization

# Spark Distro Platforms

## Spark Support Included on Big Data Platforms

- While this build process is very easy, it's even easier to have the runtime pre-built...
- Platform support also indicates stronger integration testing, supported, and integrated management tools



<http://spark.apache.org/docs/latest/hadoop-third-party-distributions.html>

# Conclusions

Pick the  
Right Tool  
for the Job



# Spark and GraphX Resources

## Spark:

- Spark Main Site (<http://spark.apache.org/docs/latest/index.html>)
- Spark Summit 2014 Conference (<http://spark-summit.org/2014>)
- Spark Configuration (<http://spark.apache.org/docs/latest/configuration.html>)
- Tuning Spark (<http://spark.apache.org/docs/latest/tuning.html>)
- Monitoring and Instrumentation (<http://spark.apache.org/docs/latest/monitoring.html>)
- Spark Overview ("Where to Go from Here" section)  
(<http://spark.apache.org/docs/latest/index.html>)

## API Docs:

- Scala  
(<http://spark.apache.org/docs/latest/api/scala/index.html#org.apache.spark.package>)
- Java (<http://spark.apache.org/docs/latest/api/java/index.html>)
- Python (<http://spark.apache.org/docs/latest/api/python/index.html>)

## GraphX:

- Graph Analytics with Graphx (<https://databricks-training.s3.amazonaws.com/graph-analytics-with-graphx.html>)
- GraphX Programming Guide (<http://spark.apache.org/docs/latest/graphx-programming-guide.html>)
- GraphX on Github (<http://amplab.github.io/graphx/>)

# Speaker Contact Information

- **InfoQ Lead Editor:**  
<http://www.infoq.com/author/Srini-Penchikala>
- **Email:** [srinipenchikala@gmail.com](mailto:srinipenchikala@gmail.com)
- **LinkedIn:** <http://linkedin.com/in/SriniPenchikala>
- **Twitter:** @srinip
- Interested in writing articles for InfoQ?
  - Send an email to: [srini@infoq.com](mailto:srini@infoq.com)

# Q&A

# Questions & Comments